

Chicago Client-Only MLS Portal + Realtor Insight Layer

Pitch packet (MVP: for-sale listings only) | Tech: Next.js + Supabase + Clerk + Vercel

What it is

- Login-only portal for the realtor's clients (VOW-style gating), not a public IDX site.
- Fast search + map: filters, neighborhood focus, commute/geo, favorites, saved searches.
- Differentiator: "Advisor" layer that combines MLS facts with the realtor's lived experience (notes/heuristics) via embeddings + retrieval.
- Guardrails: the AI cites MLS fields and retrieved agent notes; it does not invent listing facts.

Why it matters


- Clients get curated recommendations and clear tradeoffs (not endless scrolling).
- Realtor gets higher-intent conversations, fewer repetitive questions, and a consistent client experience.
- Saved searches + alerts keep clients engaged and move faster on new inventory.

MVP roadmap (4–6 weeks)

- Phase 1: Auth + gated portal + mock listings (demo).
- Phase 2: Chicago MLS ingestion (incremental sync) + normalized listings DB.
- Phase 3: Saved searches + alerts (email/SMS).
- Phase 4: Realtor knowledge base + embeddings (pgvector) + "Advisor" summaries.
- Phase 5: Conversational agent with tool-calling (searchListings, compare, explain).


Demo wireframe

Desktop: Listings + Map + Realtor Advisor



How it works (high-level)

Data is gated, auditable, and ready for AI augmentation



Scope + data/AI plan

- Scope (MVP): Chicago for-sale listings only; map + filters; favorites + saved searches.
- Data plan: ingest MLS updates into private DB (avoid live-per-query dependence).
- Compliance: login-only access, org membership, rate limits, access logging.
- AI plan: store realtor notes as documents → chunk + embed (pgvector) → retrieve for listing summaries + chat responses.

Next steps

- Confirm MLS/VOW access path (provider + credentials + display rules).
- Pick MVP neighborhoods / price bands and required filters.
- Define realtor knowledge inputs (neighborhood notes, building notes, watchouts) for the first 20–50 entries.
- Deploy the demo to Vercel and review together; then start data ingestion sprint.